

29/10/2023 - Feb

System Programming Notes

Chapter #3: Processes in Unix

→ A process is basic active entity in most OSs model.

3.1 Process Identification:

→ Each process has a process Id, parent process Id, user Id and group Id which are use for its identifications.

→ 'pid_t' is an unsigned integer type that represent a process id.

→ Synopsis:

```
#include <unistd.h>
pid_t getpid(void);
pid_t getppid(void);
```

• These functions return process id and the parent process id respectively.

- They can't return an error.
- There is no guarantee that a 'pid_t' will fit into an 'int'

②

→ To each user System administrators provide a user Id and an integral group Id.

→ The most privileged user, super-user, system administrator, or root has a user Id of 0.

- \$ sudo addgroup <group-name>
- \$ sudo usermod -aG

→ Synopsis:

```
#include <unistd.h>
```

```
gid_t getegid(void);
```

```
gid_t geteuid(void);
```

```
gid_t getgid(void);
```

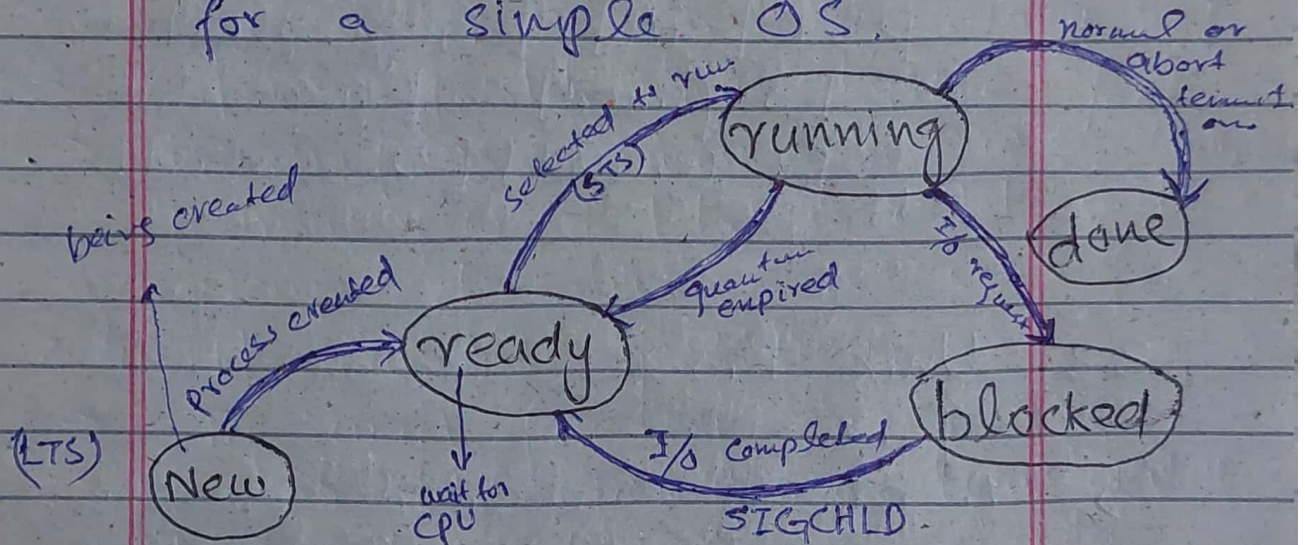
```
uid_t getuid(void);
```

- gid_t and uid_t are integral types representing group and user Ids.
- The 'getgid' and 'getuid' returns real Ids, and 'getegid' and 'geteuid' returns the effective Ids.
- None of these function returns error.

3.2 Process State:

→ The state of a process indicates its status at an time.

→ Below is state diagram for a simple OS.



→ A process perform I/O by requesting the services through system call. During the execution of system call, the OS regain control of processor and can move processes to blocked state until operation completes.

→ A context switch is moving a process from running to and replacing it with another.

(4)

→ The 'ps' utility display the information about processes.

Synopsis:

```
ps [-aA] [-G groupList]
   [-o format] ... [-p proclist]
   [-t termList] [-U userList].
```

3.3 Unix Process Creation and 'fork':

→ A process can create another process by ^{calls} fork. The caller becomes the parent and the created is called child process.

→ The fork function copies the parent memory image.

→ Synopsis

```
#include <unistd.h>
pid_t fork (void);
```

- returns 0 to child and child's pid to parent on success.

- • On failure it returns -1 and set errno. (EAGAIN).
- The output from the processes (child + parent) can appear in either order, and may not always be the same for running again and again in order.

3.4 The 'wait' function:

- The wait function causes the caller to suspend execution until a child's status becomes available or until the caller receives a signal.
- A process status becomes available after termination or after process has been stopped.
- Synopsis: `#include <sys/wait.h>`
 - `pid_t wait (int *stat_loc);`
 - `pid_t waitpid (pid_t pid, int *stat_loc, int options);`
 - If 'pid' is -1, it will

wait for any child.

- `pid` is pointer to a location for returning a status and a flag specifying options.
- If `pid > 0`, it will wait for a child whose process id is `pid`.
- If `pid < -1`, it will wait for any child in process group specified by absolute value of `pid`.
- The 'options' parameter of `waitpid` is bitwise inclusive OR of one or more flags.
- On success - the function returns `pid` of child.
- On failure it returns `-1` and set `errno`.

→ If a parent is not waiting for a child it is known as orphan zombie

→ If a parent terminates without waiting for its child, the child becomes an orphan.

- 'stat_loc' argument of 'wait' or 'waitpid' is a pointer to integer variable. If it is not NULL, these functions store return status of child in this location.
- Child return its status by `exit`, `_exit` or return from main.
- Parent only access the 8 least significant bits of child's return status.
- POSIX specifies 6 Macros for testing child's return status. Each takes the status value returned by a child to 'wait' or 'waitpid' as a parameter.
- Synopsis


```
#include <sys/wait.h>

i) WIFEXITED(int stat_val);
ii) WEXITSTATUS(    "    );
iii) WIFSIGNALED(    "    );
iv) WTERMSIG(    "    );
v) WIFSTOPPED(    "    );
vi) WSTOPSIG(    "    );
```


8

- These are designed to be used in pairs.

- (i) returns ^{non-zero} true if process terminates normally.
- (i) == true Then: (ii) evaluates to the lower 8-bits return by child through `exit`, `_exit`, `_Exit` or `return`.
- (ii) evaluates to ^{true} non-zero when child terminates because of an uncaught signal/ abnormal termination of child.
- (ii) == true Then: (iv) reason of abnormal termination, evaluates to number of signal that caused termination.
- (iv) evaluates to ^{true} non-zero value if child is currently stopped.
- (v) == true Then: (vi) evaluates to number of signal that caused the child to stop/ reason of suspension/stopping the child.

3.5 The exec Function:

→ Fork creates image of callers. The exec family functions creates a new image other than that of caller.

→ Synopsis

- ```
#include <unistd.h>
```
- extern char \*\*environ;
  - int execl(const char\* path, const char\* filename, const char\* arg[0], ..., NULL);
  - int execlp(const char\* filename, const char\* arg[0], ..., NULL);
  - int execv(const char\* path, <sup>const char\*</sup>char \*const argv[]);
  - int execvp(const char\* filename, char \*const argv[]);
  - int execlp(const char\* path, <sup>char \*const</sup>char \*const argv[], <sup>char \*const</sup>char \*const envp[]);
  - int execlp(const char\* path, char \*const argv[], char \*const envp[]);
- They don't return anything in success case



10

- These functions return -1 and set 'errno' if unsuccessful.

→ The 'exec' with 'l' family is used if we know the number of CLAs at compile time.

→ The 'exec' with 'v' family are used if we don't know the number of CLA at compile time.