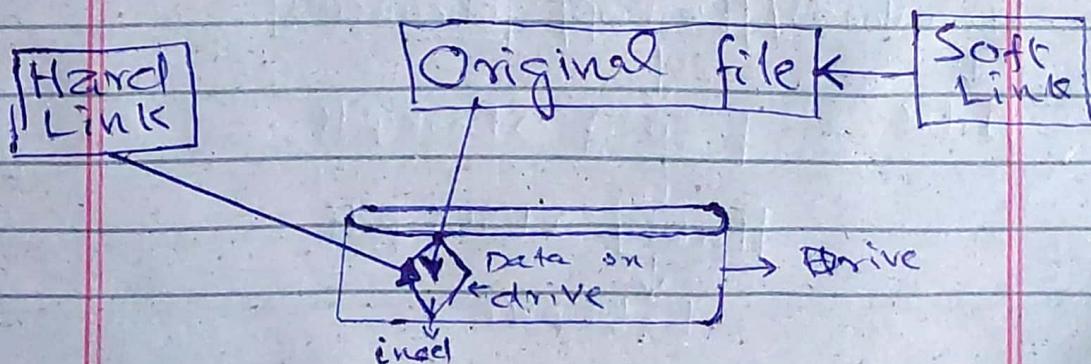


System Programming Notes

chapter # 5. Files and Directories

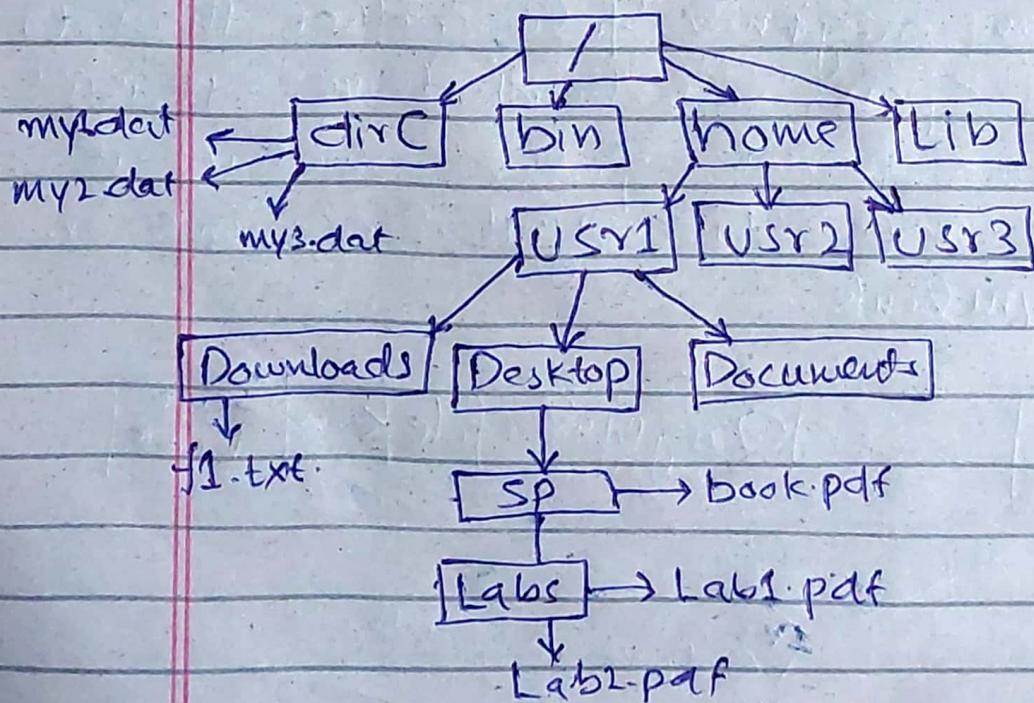
- Unix file systems are tree structured, with nodes representing files and arcs representing the contains relationship.
- Hard Link: Pointing directly to a structure containing the file location information. It is a file all its own. point to exact spot on a harddrive where inodes stores the data.
- Soft Link: points to name of the original file, rather than to a spot on hard drive. It isn't a separate file. Sometime it is refer to symbolic link or symlink.



(2)

5.1 Unix file system navigation

- Filesystem is the way in which files, directories are named and where they are placed logically for storage and retrieval. eg NTFS, FAT etc.
- A file is any computer document while a directory is special file contains information about other files and folder called directory entries direct.
- The 'root' directory of filesystem is at top of tree and everything else are branches of it.



point

mostly starts with . /
or ...

- A relative path describe the location of a file relative to current working directory. An absolute path describe the location from root directory. mostly starts with /.
- root directory has both dot and dot-dot pointing to itself.
- The 'PWD' environment variable specify the current working directory of a process. Also 'getcwd' retrieve the current working directory.

→ Synopsis:

```
#include <unistd.h>
```

```
int chdir(const char* path);
```

use to change the current working directory within a process.

~* Path : The specified path become the current working directory for calling process.

(4)

- * 0 : when successful
- * -1 : when unsuccessful and set 'errno'.

→ Synopsis

```
#include <unistd.h>
char *getcwd (char *buff,
               size_t size);
```

This function return the pathname of current working directory.

- ~ *buff : represents a user-supplied buffer for holding pathname of cwd.
- ~ size : specifies maximum length pathname that a buf can accommodate, including the trailing string terminator.
- * a pointer to buff when successful
- * returns NULL and set errno.

→ PATH_MAX is an optional POSIX constant specifying the maximum length of a pathname (including '/') for the implementation. It may or may not be defined in limits.h.

5.2 Directory Access:

→ Directory cannot be accessed with ordinary 'open', 'close', 'read' functions. Instead the required specialize function whose corresponding names end with 'dir'.

→ Synopsis:

```
#include <dirent.h>
DIR *opendir(const char* dirname);
```

This function provides a handle to a directory stream that is positioned at first entry in directory.

~ *dirname: Path to directory

* pointer to directory object when successful.

(6)

- * NULL: returns a null pointer and set errno when unsuccessful.
 - DIR type represent directory stream which is an ordered sequence of all of the directory entries in a particular order.
 - Synopsis:
`#include <dirent.h>
struct dirent *readdir(DIR *dirp);`
- The 'readdir' function reads a directory by returning a successive entries in a directory stream pointed 'dirp'. It moves the stream to next position after each call.
- ~ *dirp:
- * struct dirent *: returns a pointer to the structure containing the information about next directory entry. w
 - * NULL: also set errno.

→ readdir also return NULL to indicate the end of directory, but in this case it doesn't change errno.

→ Synopsis:

```
#include <dirent.h>
int closedir(DIR *dirp)
void rewinddir(DIR *dirp)
```

The 'closedir' closes a directory stream and the 'rewinddir' function reposition the directory stream to its beginning.

~ *dirp: correspond to an open directory stream

*^s
*^{us} 0
-1 and set errno

→ Traditionally, UNIX directory entries only contains filenames/directorynames and inode numbers.

The inode number is an index into a table containing the other information about a file.

⑧

→ We have three functions for retrieving file status info:

- fstat access file with open file descriptor.
- stat and lstat access by file name.

→ Synopsis:

```
#include <sys/stat.h>
```

- int lstat (const char* restrict path, struct stat *restrict buf);
- int stat (const char *restrict path, struct stat *restrict buf);

~ Path: name of symbolic link of a file whose status is to be returned.

~ *buf: point to user supplied buffer into which the information is stored.

*
S
*
us

O

-1 and set errno.

• `int fstat(int filedes, struct stat*buf);`

~ `filedes`: filedescriptor of open filo.

~ `buf`: as above function.

* 3
0
* -1 and set `errno`.

→ 'ctime' function takes a time and return its ^{human} & readable string format. keep in mind `ctime` returned string end with new line, we can remove the new line by using `strtok` function as:

"`strtok(ctime(time), "\n");`"

→ The `stat` contain 8 members which store time of last access, last data modification and last status change. These are stored in '`time_t st_atime, time_t st_mtime and time_t st_ctime`' respectively.

→ It has also a member named '`mode_t st_mode`' specifies access permission and type of

10

file.

- For file type Posix specifies macros (on page 197).

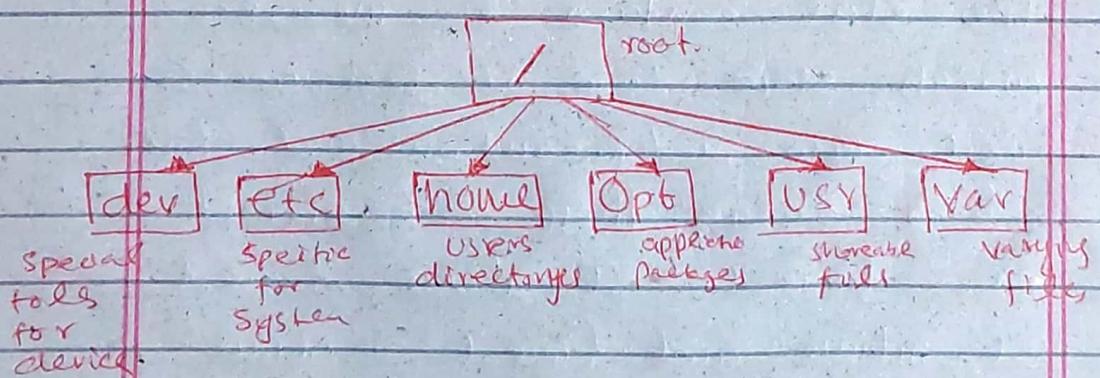
" S_ISDIR(st_mode) ? printf("d") : printf("-");

- For permission take bitwise and of permission bit and st_mode.

" S_IRUSR & ~st_mode ? printf("r") : printf("-");

5.3 Unix file system Implementation:

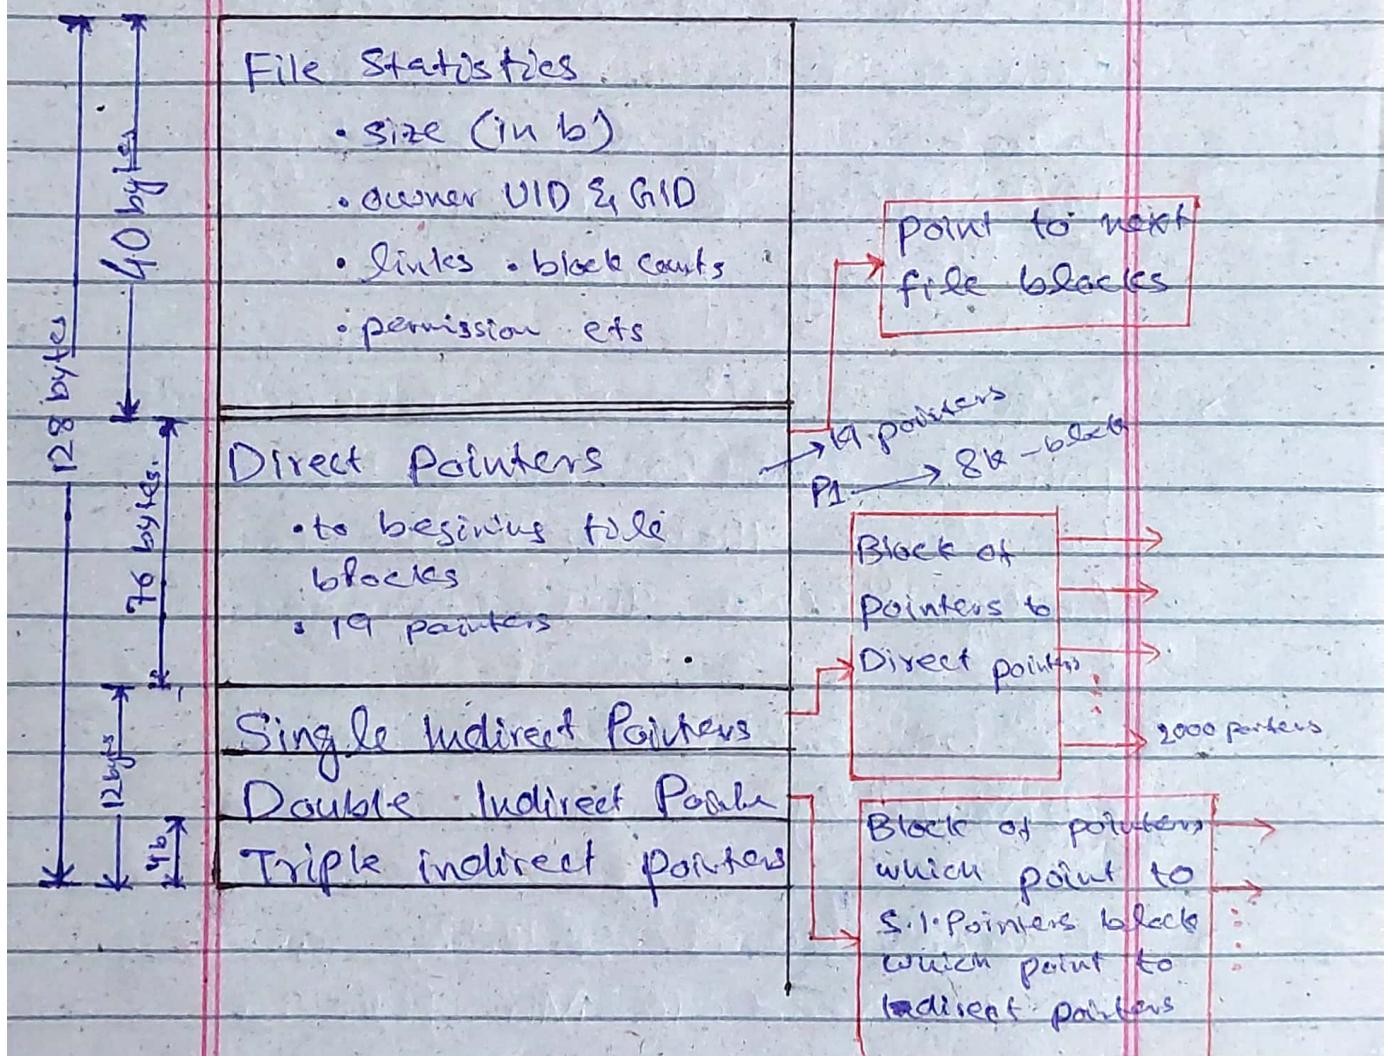
- Disk formatting divides a physical disk into regions called partitions.
- All fully qualified paths in Unix start from root directory /.



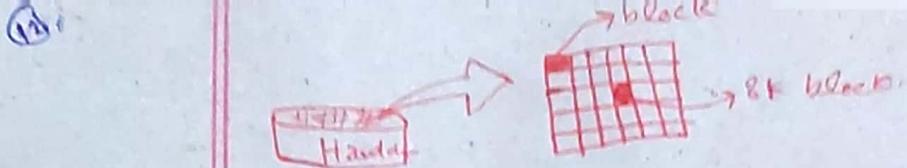
- Directory entries contain filename and a reference to fixed-length structure called inod.
- The inad contains information about

the file size, file location, owner of file, time of creation, last access and modification, permission and so on.

→ Below figure shows the indirect structure for a typical file.



- Direct pointer to first few data blocks of file.
- If file is large, the indirect



pointer is a pointer to a block of pointers that points to additional data blocks.

- Similarly double ^{indirect} pointer points to a block of single indirect pointers and so on.
- The block can mean different things. In this context a block is typically 8K bytes. The number of bytes in a block is always power of 2.
- Since $2048^3 = 2^{33}$ (T.I.P Case), pointer's ~~total~~ ^{size} would need to be longer than 4 bytes to fully address this storage.
- 32-bit addresses can access approximately 4 billion blocks.
- It takes about 20 times as long as to read 16-MB file in which the data blocks are randomly placed than one in which the data blocks are contiguous.

→ Directory implementation that only contains name and inod has following advantages:

- Changing file name or moving a file (mv command) only requires changing directory entry.
- Only one physical copy of file needs to exist on disk.
- Directory entries are of variable length because the filename is of variable length.

5.4 Hard Links and Symbolic (Soft) Links:

- A link, sometimes called a hard link, is a director entry.
- A symbolic link sometimes called a soft link is a file that stores a string used to modify the pathname when it is encountered during pathname resolution.
- See the figure on first page for better understanding.
- We can create additional

(10)

links to a file with the "ln" shell command or the link function.

→ Creation of new link adds directory entry and increments count of links pointing to inode, and by removing ('rm') the file it does vice versa.

→ Synopsis:

The link function creates a new entry for existing file. (`#include <unistd.h>`)

```
int link(const char* path1,  
        const char* path2);
```

~ Path1: path to file

~ Path2: new path of file.

*^s O

*⁽⁴⁾ -1 and set errno.

→ Synopsis:

The unlink function removes the directory entry specified by

path. If link count is 0 and no process has opened it, the unlink frees the space occupied by file.

```
#include <unistd.h>
int unlink(const char* path);
~path: path to file to be removed.
```

*^s O

*^{us} → and set errno

→ A symbolic link is a file that contains the name of another file or directory.

→ We can create a soft link by 'ln' command with '-s' option or by invoking 'symlink' function.

→ Synopsis:

```
#include <unistd.h>
```

```
int symlink(const char* path1,
           const char* path2);
```

~path1: path that is a string that will be content of link.

~path2: gives pathname of link.

(16)

*^s 0; *^{us} -1 and set errno.